

Multithreaded Algorithms

Bin Wang

School of Software
Tsinghua University

June 7, 2022

Outline

- 1 The basics of dynamic multithreading
- 2 Multithreaded matrix multiplication
- 3 Multithreaded merge sort

Overview

Parallel computers

- Highest-priced machines:
Supercomputers
- Intermediate price:
Clusters built from individual computers
- Inexpensive desktop/laptop:
Chip multiprocessors

Overview

Parallel computing models

- **shared memory** vs **distributed memory**
- **static threading**
- **concurrency platforms**

Overview

Dynamic multithreaded programming

- **nested parallelism**
- **parallel loops**
- **parallel, spawn, sync and new**
- follow naturally from the divide-and-conquer paradigm
- faithful to how parallel-computing practice is evolving

Overview

Dynamic multithreaded programming

- **Cilk**
 - MiT Cilk, from 1994
 - Cilk++, from 2006
 - Intel Cilk Plus, 2009–2017
- **Open MP**, from 1997
- **Task Parallel Library**, from 2010 (.NET Framework 4.0)
- **Task Building Blocks (Intel TBB)**, from 2006

Multithreaded Fibonacci Algorithm

The serial algorithm

$\text{FIB}(n)$

```
1  if  $n \leq 1$ 
2      return  $n$ 
3  else  $x = \text{FIB}(n - 1)$ 
4       $y = \text{FIB}(n - 2)$ 
5      return  $x + y$ 
```

Multithreaded Fibonacci Algorithm

The serial algorithm

$\text{FIB}(n)$

```
1  if  $n \leq 1$ 
2      return  $n$ 
3  else  $x = \text{FIB}(n - 1)$ 
4       $y = \text{FIB}(n - 2)$ 
5      return  $x + y$ 
```

$T(n) = \Theta(\phi^n)$, where $\phi = (1 + \sqrt{5})/2$.

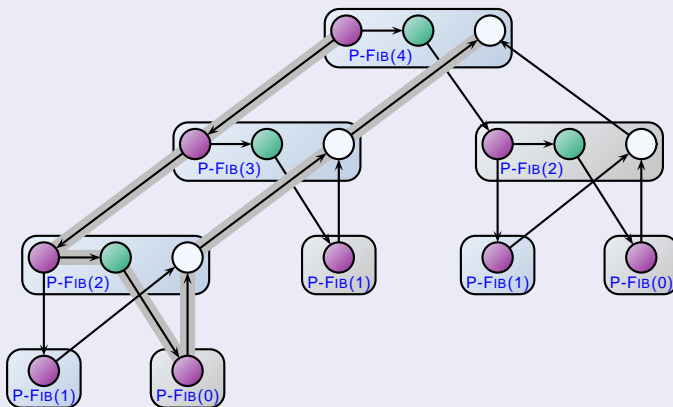
Multithreaded Fibonacci Algorithm

The multithreaded algorithm

P-FIB(n)

```
1  if  $n \leq 1$ 
2      return  $n$ 
3  else  $x = \text{spawn P-FIB}(n - 1)$ 
4       $y = \text{P-FIB}(n - 2)$ 
5      sync
6      return  $x + y$ 
```

A model for multithreaded execution



Multithreaded Fibonacci Algorithm

A model for multithreaded execution

- **Computation dag**
- **Initial strand** and **final strand**
- **Continuation edge**
- **Spawn edge**
- **Call edge**
- **Ideal parallel computer**
- **Sequentially consistent** shared memory

Performance measures

Performance measures

- The **work** of a multithreaded computation is the total time to execute the entire computation on one processor. Denote as T_1 .
- The **span** is the longest time to execute the strands along any path in the dag. Denote as T_∞ .



Performance measures

Performance measures

- *work law:*

$$T_p \geq T_1/P.$$

- *span law:*

$$T_p \geq T_\infty.$$

- *speedup :*

$$T_1/T_p.$$

Performance measures

Performance measures

- *linear speedup* :

$$T_1/T_p = \Theta(P).$$

- *perfect linear speedup* :

$$T_1/T_p = P.$$

- *parallelism* :

$$T_1/T_\infty.$$

Performance measures

Performance measures

- ***(parallel) slackness :***

$$(T_1/T_\infty)/P = T_1/(PT_\infty).$$

- As the slackness increases from 1, a good scheduler can achieve closer and closer to perfect linear speedup.

Scheduling

Scheduling

- A multithreaded scheduler must schedule the computation *on-line*.
- *Greedy schedulers* assign as many strands to processors as possible in each time step.
- If at least P strands are ready to execute during a time step, we say that the step is a *complete step*.

Scheduling

Theorem 27.1

On an ideal parallel computer with P processors, a greedy scheduler executes a multithreaded computation with work T_1 and span T_∞ in time

$$T_p \leq T_1/P + T_\infty.$$

Scheduling

Theorem 27.1

On an ideal parallel computer with P processors, a greedy scheduler executes a multithreaded computation with work T_1 and span T_∞ in time

$$T_p \leq T_1/P + T_\infty.$$

Proof.

Suppose for the purpose of contradiction that the number of complete steps is strictly greater than $\lfloor T_1/P \rfloor$.

Scheduling

Proof.

Then, the total work of the complete steps is at least

$$\begin{aligned}
 P \cdot (\lfloor T_1/P \rfloor + 1) &= P \lfloor T_1/P \rfloor + P \\
 &= T_1 - (T_1 \bmod P) + P \\
 &> T_1
 \end{aligned}$$

Contradiction! So we conclude that the number of complete steps is at most $\lfloor T_1/P \rfloor$.

Scheduling

Proof.

Now we consider an incomplete step. Let G be the dag representing the entire computation.

An incomplete step decreases the span of the unexecuted dag by 1. Hence, the number of incomplete steps is at most T_∞ .

Since each step is either complete or incomplete, the theorem follows. \square

Scheduling

Corollary 27.2

The running time T_p of any multithreaded computation scheduled by a greedy scheduler on an ideal parallel computer with P processors is within a factor of 2 of optimal.

Proof.

Since the work and span laws give us $T_p^* \geq \max(T_1/P, T_\infty)$, Theorem 27.1 implies that

Scheduling

Proof.

$$\begin{aligned}T_P &\leq T_1/P + T_\infty \\&\leq 2 \cdot \max(T_1/P, T_\infty) \\&\leq 2T_P^*\end{aligned}$$



Scheduling

Corollary 27.3

If $P \ll T_1/T_\infty$, we have $T_p \approx T_1/P$, or equivalently, a speedup of approximately P .

Proof.

If we suppose that $P \ll T_1/T_\infty$, then we also have $T_\infty \ll T_1/P$, and hence Theorem 27.1 gives us $T_p \leq T_1/P + T_\infty \approx T_1/P$.

Since $T_p \geq T_1/P$, we conclude $T_p \approx T_1/P$. \square

Analyzing multithreaded algorithms

P-FIB(n)

- $T_1(n) = T(n) = \Theta(\phi^n)$





$$\begin{aligned} T_\infty(n) &= \max(T_\infty(n-1), T_\infty(n-2)) + \Theta(1) \\ &= T_\infty(n-1) + \Theta(1) \\ &= \Theta(n) \end{aligned}$$

- The parallelism of P-FIB(n) is $T_1(n)/T_\infty(n) = \Theta(\phi^n/n)$.

Parallel loops

Multiplying matrix by and vector

MAT-VEC(A, x)

```
1   $n = A.rows$ 
2  let  $y$  be a new vector of length  $n$ 
3  parallel for  $i = 1$  to  $n$  
4       $y_i = 0$ 
5  parallel for  $i = 1$  to  $n$ 
6      for new  $j = 1$  to  $n$  
7           $y_i = y_i + a_{ij}x_j$ 
8  return  $y$ 
```

Parallel loops

Multiplying matrix by and vector

MAT-VEC-MAIN-LOOP(A, x, y, n, i, i')

```
1  if  $i == i'$ 
2      for  $j = 1$  to  $n$ 
3           $y_i = y_i + a_{ij}x_j$ 
4  else  $mid = \lfloor (i + i')/2 \rfloor$ 
5      spawn MAT-VEC-MAIN-LOOP( $A, x, y, n, i, mid$ )
6      MAT-VEC-MAIN-LOOP( $A, x, y, n, mid + 1, i'$ )
7      sync
```

Parallel loops

Multiplying matrix by and vector

- $T_1 = \Theta(n^2)$.
- $T_\infty(n) = \Theta(\lg n) + \max_{1 \leq i \leq n} \text{iter}_\infty(i)$.
- The parallel initialization loop in lines 3 – 4 has span $\Theta(\lg n)$.
- The span of the doubly nested loops in lines 5 – 7 is $\Theta(n)$.
- The parallelism is $\Theta(n^2)/\Theta(n) = \Theta(n)$.

Race conditions

Race conditions

- A multithreaded algorithm is **deterministic** if it always does the same thing on the same input.
- It is **nondeterministic** if its behavior might vary from run to run.
- A multithreaded algorithm intended to be deterministic **fails** to be, because it contains a **“determinacy race”**.

Race conditions

Example

A *determinacy race* occurs when two logically parallel instructions access the same memory location and at least one of the instructions performs a write.

RACE-EXAMPLE()

```
1  x = 0
2  parallel for i = 1 to 2
3      x = x + 1
4  print x
```

Race conditions

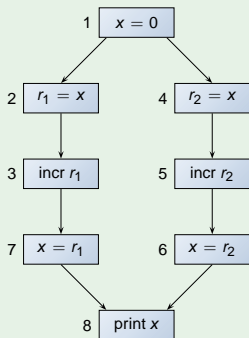
Example

When a processor increments x , the operations is composed of a sequence of instructions:

- 1 Read x from memory into one of the processor's registers.
- 2 Increment the value in the register.
- 3 Write the value in the register back into x in memory.

Race conditions

Example



step	x	r_1	r_2
1	0	—	—
2	0	0	—
3	0	1	—
4	0	1	0
5	0	1	1
6	1	1	1
7	1	1	1

Race conditions

How to cope with races?

- Mutual-exclusion locks and other methods of synchronization.
- Ensure that strands that operate are independent.
 - In a **parallel for** construct, all the iterations should be independent. Sometimes using the **new** keyword to ensure that different iterations do not operate on the same variable.
 - Between a **spawn** and the corresponding **sync**, the code of the spawned child should be independent of the code of the parent.

Race conditions

Example

MAT-VEC-WRONG(A, x)

```
1   $n = A.rows$ 
2  let  $y$  be a new vector of length  $n$ 
3  parallel for  $i = 1$  to  $n$ 
4       $y_i = 0$ 
5  parallel for  $i = 1$  to  $n$ 
6      parallel for new  $j = 1$  to  $n$ 
7           $y_i = y_i + a_{ij}x_j$ 
8  return  $y$ 
```

A chess lesson

A chess lesson

- The program was prototyped on a 32-processor computer but was ultimately to run on a supercomputer with 512 processors.
- The original version of the program had work $T_1 = 2048$ seconds and span $T_\infty = 1$ second.

A chess lesson

A chess lesson

- With optimization, the work became $T'_1 = 1024$ seconds and span became $T'_\infty = 8$ seconds.
- $T_{32} = 2048/32 + 1 = 65$ and $T'_{32} = 1024/32 + 8 = 40$
- $T_{512} = 2048/512 + 1 = 5$ and $T'_{512} = 1024/32 + 8 = 10$

Multithreaded matrix multiplication

Example

P-SQUARE-MATRIX-MULTIPLY(A, B)

```
1   $n = A.rows$ 
2  let  $C$  be a new  $n \times n$  matrix
3  parallel for  $i = 1$  to  $n$ 
4      parallel for new  $j = 1$  to  $n$ 
5           $c_{ij} = 0$ 
6          for new  $k = 1$  to  $n$ 
7               $c_{ij} = c_{ij} + a_{ik} \cdot b_{kj}$ 
8  return  $C$ 
```

Multithreaded matrix multiplication

Analyzing P-SQUARE-MATRIX-MULTIPLY

- $T_1(n) = \Theta(n^3)$
- $T_\infty(n) = \Theta(\lg n) + \Theta(\lg n) + \Theta(n) = \Theta(n)$
- The parallelism is $\Theta(n^3)/\Theta(n) = \Theta(n^2)$

A divide-and-conquer algorithm

Example

$$A = \begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix}, B = \begin{pmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{pmatrix}, C = \begin{pmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{pmatrix}.$$

$$\begin{aligned} \begin{pmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{pmatrix} &= \begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix} \begin{pmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{pmatrix} \\ &= \begin{pmatrix} A_{11}B_{11} & A_{11}B_{12} \\ A_{21}B_{11} & A_{21}B_{12} \end{pmatrix} + \begin{pmatrix} A_{12}B_{21} & A_{12}B_{22} \\ A_{21}B_{21} & A_{22}B_{22} \end{pmatrix} \end{aligned}$$

A divide-and-conquer algorithm

Example

P-MATRIX-MULTIPLY-RE(C, A, B)

1 $n = A.rows$

2 **if** $n == 1$

3 $c_{11} = a_{11}b_{11}$

4 **else** let T be a new $n \times n$ matrix

5 partition A, B, C and T into $n/2 \times n/2$

submatrices: $A_{11}, A_{12}, A_{21}, A_{22};$

$B_{11}, B_{12}, B_{21}, B_{22}; C_{11}, C_{12}, C_{21}, C_{22};$

and $T_{11}, T_{12}, T_{21}, T_{22};$ respectively

6 **spawn** P-MATRIX-MULTIPLY-RE(C_{11}, A_{11}, B_{11})

7 **spawn** P-MATRIX-MULTIPLY-RE(C_{12}, A_{11}, B_{12})

A divide-and-conquer algorithm

Example

```

8      spawn P-MATRIX-MULTIPLY-RE( $C_{21}, A_{21}, B_{11}$ )
9      spawn P-MATRIX-MULTIPLY-RE( $C_{22}, A_{21}, B_{12}$ )
10     spawn P-MATRIX-MULTIPLY-RE( $T_{11}, A_{12}, B_{21}$ )
11     spawn P-MATRIX-MULTIPLY-RE( $T_{12}, A_{12}, B_{22}$ )
12     spawn P-MATRIX-MULTIPLY-RE( $T_{21}, A_{22}, B_{21}$ )
13     P-MATRIX-MULTIPLY-RE( $T_{22}, A_{22}, B_{22}$ )
14     sync
15     parallel for  $i = 1$  to  $n$ 
16         parallel new for  $j = 1$  to  $n$ 
17              $C_{ij} = C_{ij} + t_{ij}$ 

```


A divide-and-conquer algorithm

Analyzing the running time

- The work

$$M_1(n) = 8M_1(n/2) + \Theta(n^2) = \Theta(n^3)$$

- The span

$$M_\infty(n) = M_\infty(n/2) + \Theta(\lg n) = \Theta(\lg^2 n)$$

- Its parallelism is

$$M_1(n)/M_\infty(n) = \Theta(n^3/\lg^2 n)$$

Multithreading Strassen's method

Multithreading Strassen's method

- [1] Divide the input matrices A and B and output matrix C into $n/2 \times n/2$ submatrices. This step takes $\Theta(1)$ work and span by index calculation.
- [2] Create 10 matrices S_1, S_2, \dots, S_{10} , each of which is $n/2 \times n/2$ and is the sum or difference of two matrices created in step 1. We can create all 10 matrices with $\Theta(n^2)$ work and $\Theta(\lg n)$ span by using doubly nested **parallel for** loops.

Multithreading Strassen's method

Multithreading Strassen's method

- [3] Using the submatrices created in step 1 and the 10 matrices created in step 2, recursively spawn the computation of seven $n/2 \times n/2$ matrix products P_1, P_2, \dots, P_7 .
- [4] Compute the desired submatrices $C_{11}, C_{12}, C_{21}, C_{22}$ of the result matrix C by adding and subtracting various combinations of the P_i matrices, once again using doubly nested **parallel for** loops.

A naive multithreaded merge sort

Example

$\text{MERGE-SORT}'(A, p, r)$

```
1  if  $p < r$ 
2       $q = \lfloor (p + r) / 2 \rfloor$ 
3      spawn  $\text{MERGE-SORT}'(A, p, q)$ 
4       $\text{MERGE-SORT}'(A, q + 1, r)$ 
5      sync
6       $\text{MERGE}(A, p, q, r)$ 
```

A naive multithreaded merge sort

Analyzing the running time

- The work

$$MS'_1(n) = 2MS'_1(n/2) + \Theta(n) = \Theta(n \lg n)$$

- The span

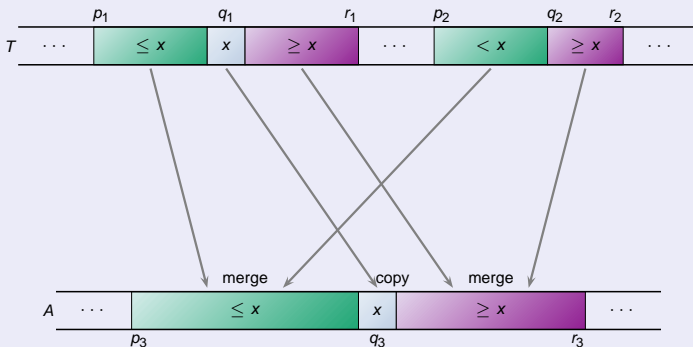
$$MS'_\infty(n) = MS'_\infty(n/2) + \Theta(n) = \Theta(n)$$

- Its parallelism is

$$MS'_1(n)/MS'_\infty(n) = \Theta(\lg n)$$

Multithreaded merging

Parallel Merging



Multithreaded merging

Parallel Merging

P-MERGE($T, p_1, r_1, p_2, r_2, A, p_3$)

```
1   $n_1 = r_1 - p_1 + 1$ 
2   $n_2 = r_2 - p_2 + 1$ 
3  if  $n_1 < n_2$ 
4      exchange  $p_1$  with  $p_2$ 
5      exchange  $r_1$  with  $r_2$ 
6      exchange  $n_1$  with  $n_2$ 
7  if  $n_1 == 0$ 
8      return
```

Multithreaded merging

Parallel Merging

```

9  else  $q_1 = \lfloor (p_1 + r_1)/2 \rfloor$ 
10       $q_2 = \text{BINARY-SEARCH}(T[q_1], T, p_2, r_2)$ 
11       $q_3 = p_3 + (q_1 - p_1) + (q_2 - p_2)$ 
12       $A[q_3] = T[q_1]$ 
13      spawn  $\text{P-MERGE}(T, p_1, q_1 - 1, p_2, q_2 - 1, A, p_3)$ 
14       $\text{P-MERGE}(T, q_1 + 1, r_1, q_2, r_2, A, q_3 + 1)$ 
15      sync
```


Multithreaded merging

Parallel Merging

BINARY-SEARCH(x, T, p, r)

```
1   $low = p$ 
2   $high = \max(p, r + 1)$ 
3  while  $low < high$ 
4       $mid = \lfloor (low + high) / 2 \rfloor$ 
5      if  $x \leq T[mid]$ 
6           $high = mid$ 
7      else  $low = mid + 1$ 
8  return  $high$ 
```

Multithreaded merging

Analysis of multithreaded merging

$$\begin{aligned}\lfloor n_1/2 \rfloor + n_2 &\leq n_1/2 + n_2/2 + n_2/2 \\ &= (n_1 + n_2)/2 + n_2/2 \\ &\leq n/2 + n/4 \\ &= 3n/4\end{aligned}$$

$$PM_{\infty}(n) = PM_{\infty}(3n/4) + \Theta(\lg n) = \Theta(\lg^2 n)$$

Multithreaded merging

Analysis of multithreaded merging

- $PM_1(n) = \Omega(n)$
- $PM_1(n) = PM_1(\alpha n) + PM_1((1 - \alpha)n) + O(\lg n)$
- We prove that $PM_1 = O(n)$ via the substitution method.
- The parallelism of P-MERGE is
 $PM_1(n)/PM_\infty(n) = \Theta(n/\lg^2 n)$

Multithreaded merge sort

Multithreaded merge sort

P-MERGE-SORT(A, p, r, B, s)

```

1   $n = r - p + 1$ 
2  if  $n == 1$ 
3       $B[s] = A[p]$ 
4  else let  $T[1 \dots n]$  be a new array
5       $q = \lfloor (p + r) / 2 \rfloor$ 
6       $q' = q - p + 1$ 
7      spawn P-MERGE-SORT( $A, p, q, T, 1$ )
8      P-MERGE-SORT( $A, q + 1, r, T, q' + 1$ )
9      sync
10     P-MERGE( $T, 1, q', q' + 1, n, B, s$ )
```

Multithreaded merge sort

Analysis of multithreaded merge sort

- The work

$$PMS_1(n) = 2PMS_1(n/2) + \Theta(n) = \Theta(n \lg n)$$

- The span

$$\begin{aligned} PMS_\infty(n) &= PMS_\infty(n/2) + \Theta(\lg^2 n) \\ &= \Theta(\lg^3 n) \end{aligned}$$

- Its parallelism is

$$PMS_1(n)/PMS_\infty(n) = \Theta(n/\lg^2 n)$$